



# NON-BLOCKING IO TECHNIQUES

Idle hands are the Devil's  
playground.

# STRATEGIES

Blocking Sockets  
(One Thread)

Non-Blocking  
Sockets

Blocking Sockets  
(Multiple Threads)

Event-driven IO

`select()`

Advanced  
Techniques

# BLOCKING SOCKETS (ONE THREAD)

*"Just live with it!"*

Pros:

- Simple

Cons:

- Non-interactive



# BLOCKING SOCKETS (MULTIPLE THREADS)

*“One thread blocks, the other runs free”*

Pros:

- Efficient
- The socket code is easy to write

Cons:

- Threading is hard



# SELECT()

*“Which of these sockets is awaiting action?”*

Pros:

- Handles large numbers of sockets well

Cons:

- Complex API
- Not very efficient for one or two sockets



# NON-BLOCKING SOCKETS

*"Are we there yet? Are we there yet? Are we there yet?"*

Pros:

- Simple

Cons:

- Polling isn't very efficient...
- ...unless your program is already structured that way





# EVENT-DRIVEN IO

*“Handle mouse, handle keyboard, handle network...”*

Pros:

- Simple
- Efficient

Cons:

- Platform-specific



# ADVANCED TECHNIQUES

*“Completion Ports and Overlapped IO”*

Pros:

- Maximally efficient for large numbers of sockets

Cons:

- Maximally complex





# STRATEGIES

Blocking Sockets  
(One Thread)

Blocking Sockets  
(Multiple Threads)

`select()`

Non-Blocking  
Sockets

Event-driven IO

Advanced  
Techniques

# NON-BLOCKING SOCKETS

1. Create the socket as usual
2. Tell the OS to set the “non-blocking IO” flag
3. Watch for WSAEWOULDBLOCK errors

```
int mode = 1;
int err = ioctlsocket(sock,
                      FIONBIO, &mode);

if (err != NO_ERROR)
    printf("failure!");
```

# NON-BLOCKING SOCKETS: RECV() AND SEND()

1. Call as normal
2. Positive result = success
3. Negative result = error
  1. If WSAEWOULDBLOCK, call again!
4. For TCP `recv()` *only*, 0 result = socket closed

```
int bytes = recv(sock, buffer,
                  maxBytes, 0);

if (bytes == SOCKET_ERROR)
    return -1;
else
    return bytes;
```

# NON-BLOCKING SOCKETS: ACCEPT()

1. Call as normal
2. Valid socket result = success
3. INVALID\_SOCKET  
result = error
  1. If WSAEWOULDBLOCK,  
call again!

```
SOCKET result = accept(sock,  
                        incoming, &size);  
  
if (result == INVALID_SOCKET)  
    return -1;  
  
return result;
```

# NON-BLOCKING SOCKETS: CONNECT()

1. Call as normal
2. Non-error result = success
  1. But that won't happen
3. "Errors":
  1. WSAEWOULDBLOCK = try again!
  2. WSAEINVAL = ready to use!
  3. WSAEALREADY = ready to use!

```
int CheckConnect(SOCKET sock, sockaddr_in*
address)
{
    if (connect(sock, (sockaddr*)address,
        sizeof(sockaddr_in)) == SOCKET_ERROR)
    {
        int error = WSAGetLastError();
        if (error != WSAEINVAL && error !=
            WSAEALREADY)
            return 0;
    }
    return 1;
}
```

# SELECT()

- `int select(  
 int nfd,  
 fd_set* readfds,  
 fd_set* writefds,  
 fd_set* exceptfds,  
 const struct timeval* timeout  
);`

```
FD_ZERO(&fds);  
FD_CLR(index, &fds);  
FD_SET(index, &fds);  
FD_ISSET(index, &fds);
```

- Still have to use non-blocking sockets, because Linux is naughty!
- `fd_set` is both input and output, so you have to reset it between calls.
- `timeout` is also both input and output on Linux, so you have to reset it between calls too.



# SUMMARY

- Many techniques available—including just not bothering!—but the best ones in the context of this class are either multithreading, non-blocking sockets or `select()`.
- Multithreading is hard. Don't do it unless you're doing it already.
- Non-blocking sockets are mostly easy.
- `select()` is only worthwhile if you need it for non-blocking stdio on POSIX.

Non-blocking sockets:

- Use `ioctlsocket()` to set `FIONBIO` flag to 1
- Treat `WSAEWOULDBLOCK` “error” as “try again”
- `Connect()` is a special case—success is represented by “errors” `WSAEINVAL` and/or `WSAEALREADY`

