

FUNDAMENTALS OF STATE REPLICATION



Definitions and Principles

First things first

1


“

A ***distributed system*** is one in which
two or more processes work together
to perform a task.



“


*The total information processed by the system is known as its **state**.*






“

*Ensuring that all processes in the system have
the same view of the state is known as*
replication
or (less precisely)
synchronization







TWO CONFLICTING GOALS

Consistency

A distributed system is **consistent** when all its processes agree on its state.

Availability

A distributed system is **available** when all its processes are allowed to modify the state.

**WE CAN'T HAVE BOTH.
ALL WE CAN DO IS PICK WHERE WE FALL BETWEEN THEM:**



TWO COMPETING CHALLENGES

Uncertainty


When a distributed system is **consistent**, we have eliminated its uncertainty.

Time

When a distributed system is **available**, it is functioning as if time were not a factor.

**WE CAN'T AVOID BOTH.
ALL WE CAN DO IS PICK WHERE WE FALL BETWEEN THEM:**






Standard Approaches


Don't reinvent the wheel

2


Local Multiplayer




Lockstep Synchrony



Lockstep Synchrony




Lockstep Synchrony




Gotchas

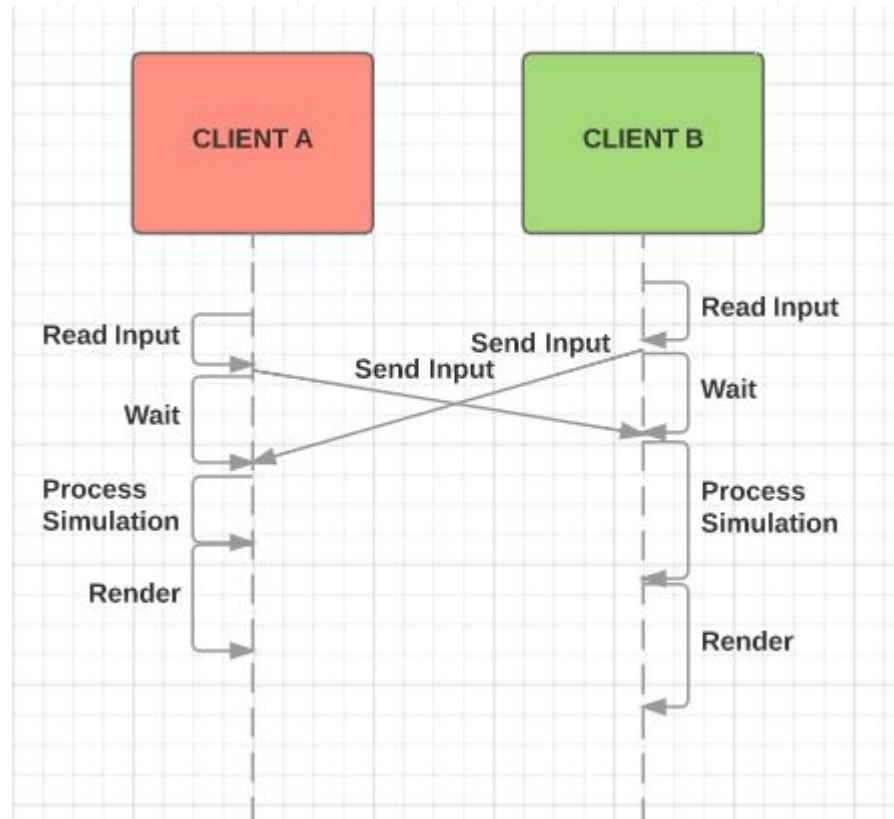
- Simulation must be completely deterministic... so be careful with the random number generator
- Object references can't use pointers





**THIS IS GOOD
ENOUGH***





Lockstep Synchrony Sequence Diagram

Lag

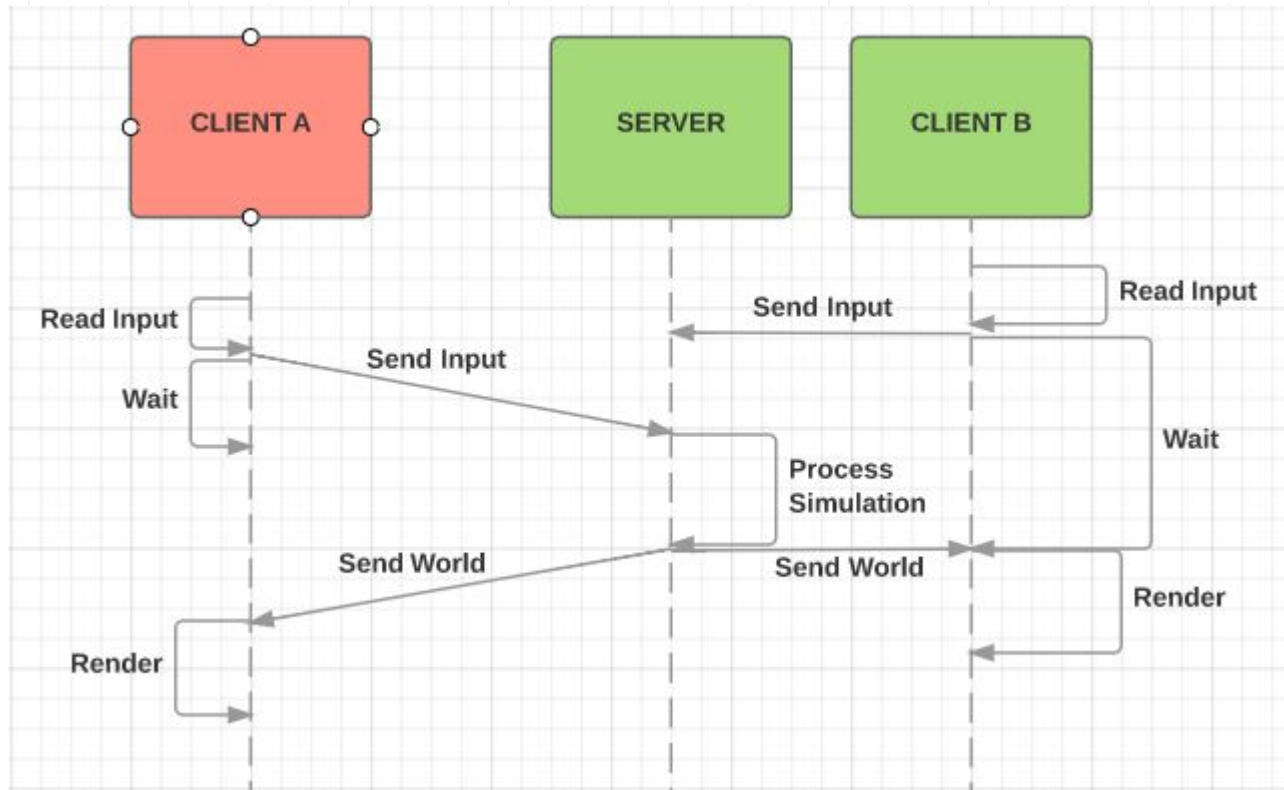
- Direct control tasks: Noticed at 50ms, annoying at 100ms, fatal at 350ms+
- Indirect control tasks: Noticed at 250ms, annoying at 1,000ms, fatal at 5,000ms+
- “Ping time” is round trip; one way should *generally* be half that



Lockstep Synchrony

- Lag is one-way trip time (or roughly $\frac{1}{2}$ ping)
- Clients are coupled; sim runs at speed of slowest client
- Rejoining difficult to impossible
- Consistency > Availability when operating normally
- Availability > Consistency when partitioned



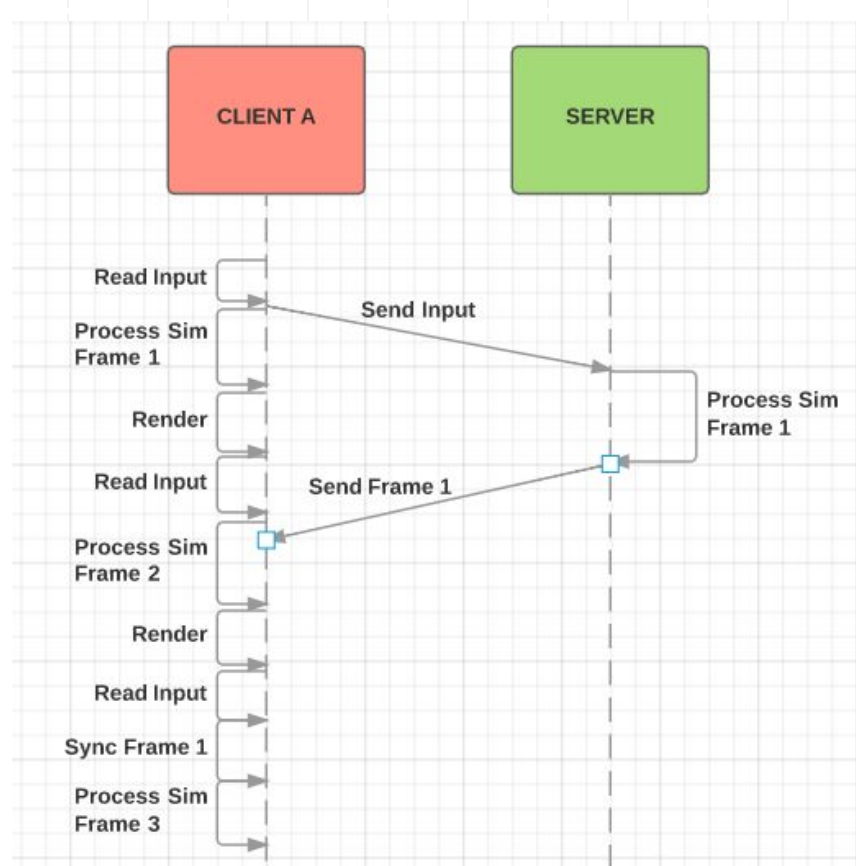


Dumb Client Synchrony Sequence Diagram

Dumb Client Synchrony

- ⦿ Lag is full ping time
- ⦿ Clients are decoupled...
- ⦿ ...but this is potentially unfair
- ⦿ Rejoining easy
- ⦿ Consistency > Availability when operating normally
- ⦿ Complete unavailable when partitioned





Optimistic Synchrony Sequence Diagram

Optimistic Synchrony

- Effectively no lag on the player's action
- But the player is always acting in the future
- Availability > Consistency in normal operation...
- ...in fact *never* consistent


Failure cases:

- Sometimes the player's action is rejected
- The player's always aiming behind targets




Optimistic Synchrony: Player Snap-Back

1. Remember player's sequence of moves
2. Snap back to divergent frame
3. Replay sequence of moves from now-corrected state
4. Display this to the player in some suitable manner



Optimistic Synchrony: Target Prediction

1. Track momentum for all remote entities
2. Project remote entities forward in time from last-known frame to current frame
3. When errors arise, they will be in the past. Rewind, integrate correct state, and project forward to the present again
4. Display this to the player in some suitable manner
5. For perceived fairness, server must evaluate critical actions from the known historical perspective of the player



Where's the simulation?

Just on the server:

- Simpler to implement
- Highly secure
- Lag on every action
- Very dependent on simulation frame rate

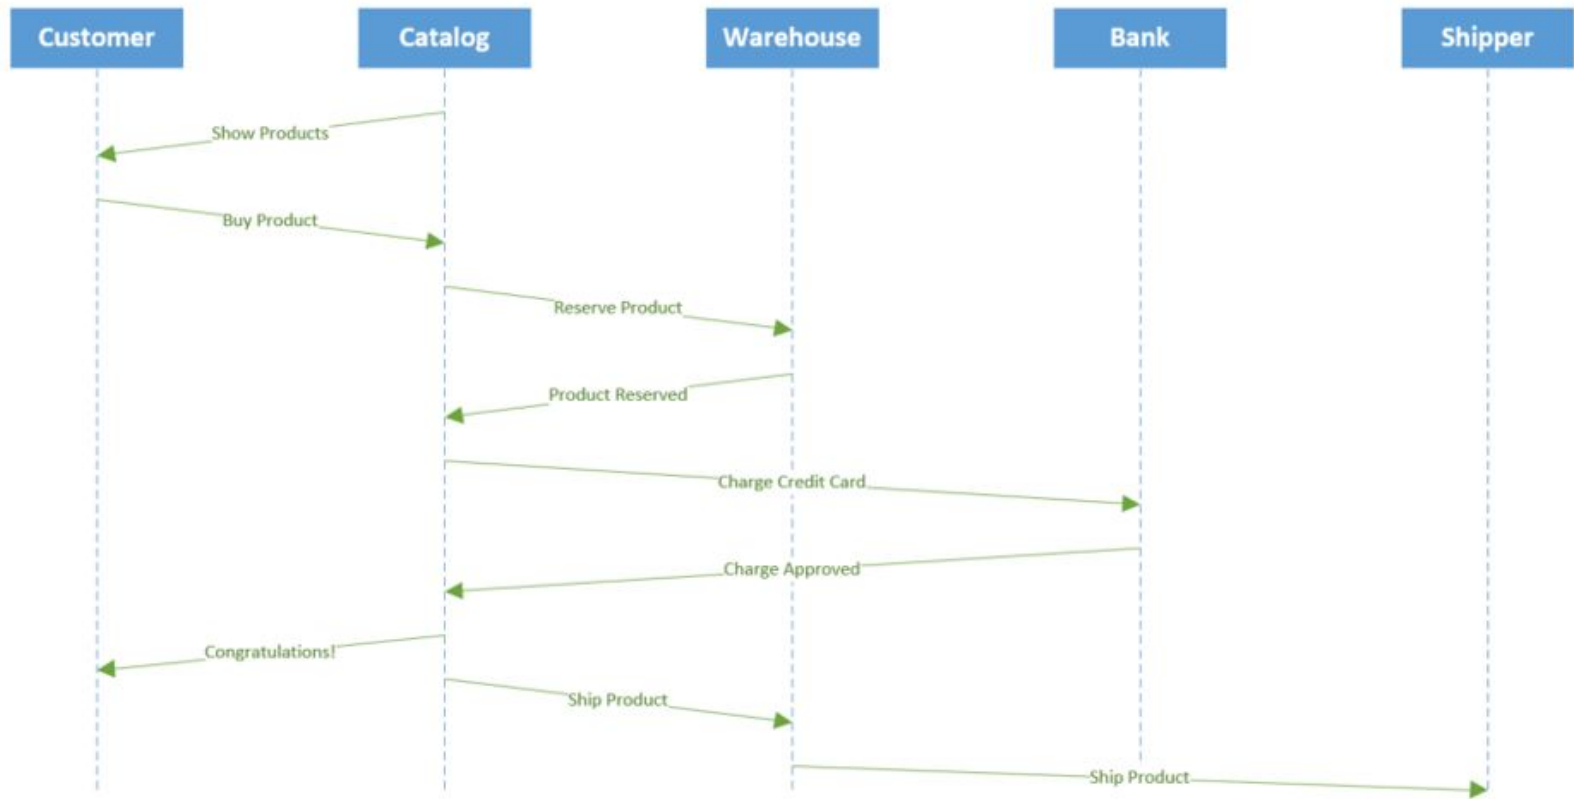
Call of Duty and World of Warcraft work this way

On both client and server:

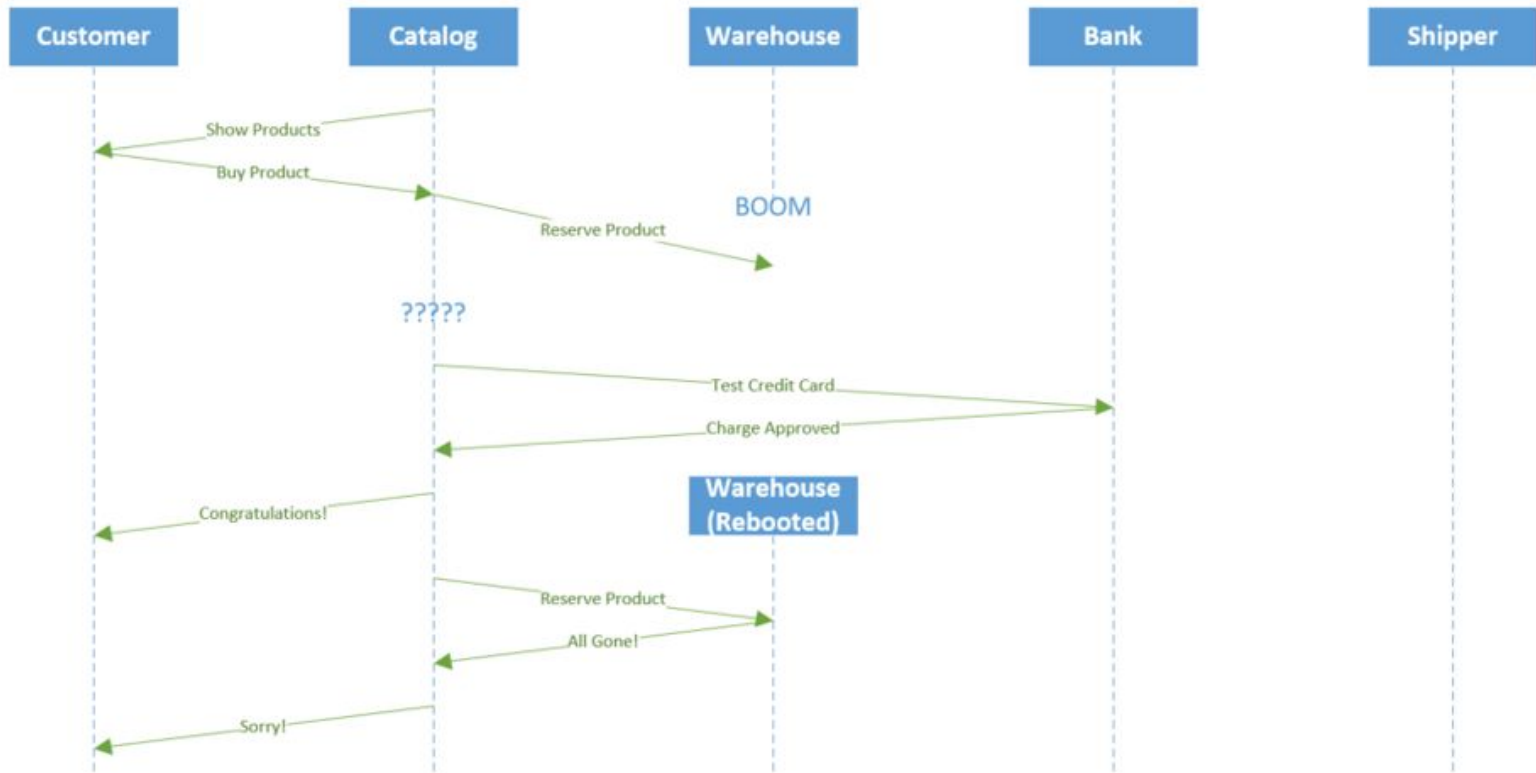
- Player's actions are lag-free
- Runs well at lower frame rates
- Physics interactions between controlled objects (e.g. players) still subject to full lag

HALO, Destiny and Assassin's Creed work this way






Eventual-Correctness Synchrony Sequence Diagram



Eventual-Correctness Synchrony Sequence Diagram




Reality Bites

If this is so easy, why doesn't everybody do it?

3

What Can Possibly Go Wrong

- Lag
- Jitter
- Packet loss
- Limited bandwidth



Lag

Lag interferes with the player's sense of control.

It comes from:

- Processing delay (number of hops)
- Queuing delay (congestion)
- Transmission delay (bandwidth)
- Propagation delay (distance)

Handle using the previously-mentioned techniques.




Jitter

Jitter is variation in lag. It disrupts our lag-tolerance schemes and makes the simulation choppy.

It comes from:

- Variation in routing (congestion and outages)
- Queuing (congestion)
- Packet loss

Handle by deepening lag compensation; only restore “normal” levels gradually.



Packet Loss

Packet loss reduces bandwidth and introduces jitter.

It comes from:

- Hardware faults
- Analog noise and radio interference
- Congestion

TCP handles packet loss automatically (if poorly); for UDP we have to develop special schemes.




Limited Bandwidth


Limited bandwidth effectively introduces lag.

It comes from:

- Poor infrastructure investment
- Lack of competition for “last mile” service
- Bad router design decisions

Handle it through compression or delaying less-important updates. A good rule of thumb is to fit within 8KB/sec.





Easy Mode

Quick and dirty networking for GAM projects

4

Easy Mode

1. Use “dumb client” networking
2. Use TCP
3. Don't run on wifi
4. If needed for your game, project simulation forward one frame while waiting for the next frame to arrive. Patch over errors using simple linear interpolation.






Summary

And in conclusion...

5

Summary

- A distributed system is one in which state is modified by more than one process...
- ...which means processes will inevitably disagree on what that state actually is.
- Either halt until disagreement is resolved (consistency over availability)...
- ...or keep on working and cope with the disagreement (availability over consistency)



Networking Models

Lockstep

All clients send input to each other, then advance simulation one frame together.


Dumb Client

All clients send input to server, which advances simulation one frame itself and sends results back to clients.

Optimistic

All clients send input to server *and* advance the simulation locally, patching up discrepancies as they arise.





CREDITS

Special thanks to all the people who made and released these awesome resources for free:

- Presentation template by [SlidesCarnival](#)
- Photographs by [Unsplash](#)



PRESENTATION DESIGN

This presentation uses the following typographies and colors:

- Titles: **Oswald**
- Body copy: **Source Sans Pro**

You can download the fonts on this page:

<https://www.google.com/fonts#UsePlace:use/Collection:Source+Sans+Pro:400,700|Oswald:400,700>

Click on the “arrow button” that appears on the top right

Sky blue **#00cef6** / Bright green **#aff000** / Blue **#3c78d8** /  **#28324a**

You don't need to keep this slide in your presentation. It's only here to serve you as a design guide if you need to create new slides or download the fonts to edit the presentation in PowerPoint®





SlidesCarnival icons are editable shapes.

This means that you can:

- Resize them without losing quality.
- Change fill color and opacity.
- Change line color, width and style.

Isn't that nice? :)

Examples:

