

MAT 345 - PROJECT #5
due Wednesday, December 12, 2018 at 12:00PM.

OBJECTIVE: In this project, you will implement a neural network.

GRADING: The assignment is worth 5% of your course grade.

INSTRUCTIONS: Students will work individually on this project, but they may ask questions and clarification from classmates and the instructor. Students must submit their projects on Moodle.

SUBMIT THE FOLLOWING: A copy of your code and a report.

PROJECT: In this project, you will build a neural network for digit recognition. Download the MNIST data set from <http://yann.lecun.com/exdb/mnist/>. There should be four files: the *training set* contains 60000 examples, and the *test set* contains 10000 examples.

0. Map output values into vectors. For example, if $y = 3$, let $\mathbf{y} = [0, 0, 1, 0, 0, 0, 0, 0, 0, 0]^T$, with a 1 in the 3rd row of a 10-dimensional vector. Note that $y = 0$ is mapped to $\mathbf{y} = [0, 0, 0, 0, 0, 0, 0, 0, 0, 1]^T$, with 1 in the 10th row.
- I. Set up the network architecture: we will use 3 layers
 - (a) The input layer will have $28 \times 28 = 784$ units since we are using images of size 28x28
 - (b) The output layer will have 10 units (1 for digit 1, \dots , 9 for digit 9, 10 for digit 0)
 - (c) The hidden layer will have s_2 units. Your code should work for different values of s_2 .
- II. Train the neural network using $\mathcal{D}_1 = \text{training set}$:

Step 1. Implement the activation function, for which we use the sigmoid $\theta(x)$ discussed in class

$$\theta(x) = \frac{e^x}{1 + e^x} = \frac{1}{1 + e^{-x}}.$$

Recall that implementation for a vector $\theta(\mathbf{z})$ is done coordinate-wise.

Step 2. Initialize the weights in $W^{(1)}$ and $W^{(2)}$ with random values in $[-\epsilon, \epsilon]$, for some small ϵ .

Note that the dimensions of the weight matrices are:

$$\begin{array}{cc} W^{(1)} & W^{(2)} \\ s_2 \times 785 & 10 \times (s_2 + 1) \end{array}$$

Step 3. For each data point $(\mathbf{x}, \mathbf{y}) \in \mathcal{D}_1$:

(1). Implement the **Feed Forward** algorithm:

(a) Let $\mathbf{a}^{(1)} = \begin{bmatrix} 1 \\ \mathbf{x} \end{bmatrix}$, that is, add the bias term $\mathbf{a}_0^{(1)} = 1$.

(b) Let $\mathbf{z}^{(2)} = W^{(1)}\mathbf{a}^{(1)}$ (it has dimension $s_2 \times 1$)

(c) Let $\mathbf{a}^{(2)} = \theta(\mathbf{z}^{(2)})$ and add $\mathbf{a}_0^{(2)} = 1$, the bias term

(d) Let $\mathbf{z}^{(3)} = W^{(2)}\mathbf{a}^{(2)}$ (it has dimension 10×1)

(e) Let $\mathbf{a}^{(3)} = \theta(\mathbf{z}^{(3)})$

(2). Implement the **Back Propagation** algorithm

(a) Let $\delta^{(3)} = \mathbf{a}^{(3)} - \mathbf{y}$ (it has dimension 10×1)

(b) Let $\tilde{W}^{(2)}$ be the matrix $W^{(2)}$ with the column of ones removed, so it has size $10 \times s_2$.

Let $\delta^{(2)} = [\tilde{W}^{(2)}]^T \delta^{(3)} \odot \theta'(\mathbf{z}^{(2)})$ (it has dimension $s_2 \times 1$).

(c) Compute the set of partial derivatives with respect to the weights, for $k = 1, 2$:

$$\begin{aligned} gW^{(1)}(\mathbf{x}, \mathbf{y}) &= \delta^{(2)}[\mathbf{a}^{(1)}]^T \rightarrow \text{size } s_2 \times 785. \\ gW^{(2)}(\mathbf{x}, \mathbf{y}) &= \delta^{(3)}[\mathbf{a}^{(2)}]^T \rightarrow \text{size } 10 \times (s_2 + 1). \end{aligned}$$

Notes: Recall that \odot stands for coordinate-wise multiplication, i.e, if $\mathbf{u} = [u_1, u_2, u_3]$ and $\mathbf{v} = [v_1, v_2, v_3]$, then

$$\mathbf{u} \odot \mathbf{v} = [u_1 v_1, u_2 v_2, u_3 v_3]$$

Also, we have derived in class that

$$\theta'(\mathbf{z}^{(k)}) = \mathbf{a}^{(k)} \odot (\mathbf{1} - \mathbf{a}^{(k)}).$$

Step 4. Compute the gradients for the *training set* \mathcal{D}_1 : for $k = 1, 2$

$$\text{grad}W^{(k)} = \frac{1}{|\mathcal{D}_1|} \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{D}_1} gW^{(k)}(\mathbf{x}, \mathbf{y}).$$

Step 5. Run the **Gradient Descent** algorithm to find the weights that minimize the cost function.

(a) Weights were initialized in Step 2.

(b) Choose η .

(c) Update rule: for $k = 1, 2$

$$W_{t+1}^{(k)} = W_t^{(k)} - \eta \cdot \text{grad}W_t^{(k)}$$

Note that the matrices and gradients can be "unrolled" as discussed in class, so one can work with vectors instead of matrices, but that is not necessary.

- III. Test the network using $\mathcal{D}_2 = \text{test set}$: compute the accuracy for this network
- (a) Use the Feed Forward part of your program to predict the output for data points in $\mathcal{D}_2 = \text{test set}$. Make sure you predict by choosing the label with the highest activation value of at least 0.5.
 - (b) Count how many data points from \mathcal{D}_2 are accurately predicted.
- IV. In your Report, include:
- (a) Your name
 - (b) The programming language you used for the project
 - (c) A discussion on which values of η lead to a reasonable performance in the gradient descent algorithm (for example, try values between 0.01 and 5.)
 - (d) Train the network with $s_2 = 30$ and output the resulting weights $W^{(1)}$, $W^{(2)}$. Test your network and output the accuracy.
 - (e) Try different number of units in the hidden layer s_2 and output the accuracy rate for each case (try a variety of sizes for s_2 , such that 30, 100, 300 etc.) Do more units in the hidden layer lead to better accuracy?
 - (f) Include any additional information, such as if you are using regularization, if you are using stochastic gradient descent rather than gradient descent, if you are doing a gradient check for back propagation, etc.